

Twits: perusing GNIP Twitter data



By Dodge Coates

Lots and lots of Tweets

Lots and lots of Tweets

- ~2 terabytes in size, uncompressed

Lots of data

Lots of data

- ~5,500 json files

Lots of data

- ~5,500 json files
- ~2 terabytes in size, uncompressed

Lots of data

- ~5,500 json files
- ~2 terabytes in size, uncompressed'
- This was reduced to ~250 gigabytes of csv

Lots of data

- ~5,500 json files
- ~2 terabytes in size, uncompressed
- This was reduced to ~250 gigabytes of csv
- No schema available for the dataset

Gnip Enterprise Access to Twitter Data

- Provide 1/10 of the twitter firehose
- Accessed for roughly 2 ½ weeks, between late January and mid February, 2012



Data Sample

```
"id_str": "163372337134182401",
"in_reply_to_status_id": null,
"created_at": "Sat Jan 28 21:26:15 +0000 2012",
. . .
"user":
{
  "id_str": "44153313",
  "favourites_count": 25,
  "contributors_enabled": false,
  . . .
  "lang": "en",
  "utc_offset": -21600,
  "profile_sidebar_border_color": "BDDCAD",
  "followers_count": 349,
  "url": "http://kingsofkauffman.com/",
  "profile_image_url": "http://a1.twimg.com/profile_images/1442806914/My_Photo__normal.jpg"
},
"retweet_count": 3,
"favorited": false,
"id": 163372337134182401,
"entities":
{ "hashtags": [
  {
    "text": "Mizzou",
    . . .
  }
],
"retweeted_status":
{
  "id_str": "163369042332229632",
  "place": null, "geo": null,
  "text": "Funny line from #Mizzou's Kim English about DGB: \"We were 100 percent focused on toda",
  "in_reply_to_status_id_str": null,
  "coordinates": null,
  "user":
  {
    "id_str": "168902884",
    "favourites_count": 0,
    "profile_use_background_image": true,
    "screen_name": "tpkcstar",
    . . .
    "favorited": false,
    "id": 163369042332229632,
  }
}
}
}
```

Goals

My Goals

1.

2.

3.



Goals

1. Exploring the data

Goals

1.Exploring the data

2.Developing an interesting prediction

Goals

- 1.Exploring the data
- 2.Developing an interesting prediction
- 3.Get a feel for working with a large dataset

The Data

- A lot to process!

Storage

- SQL

Storage

- SQL
- Distributed file systems (Hadoop)
 - Hardware is too expensive
 - API is too time-consuming

Storage

- SQL
- Distributed file systems (Hadoop)
 - Hardware is too expensive
 - API is too time-consuming
- Settled on sampling

Preprocessing



Preprocessing

- Takes about 20-30 hours on my desktop, depending on the complexity of the preprocessing

Preprocessing

- Takes about 20-30 hours on my desktop, depending on the complexity of the preprocessing
- Produces ~5,500 csv files

Preprocessing

- Takes about 20-30 hours on my desktop, depending on the complexity of the preprocessing
- Produces ~5,500 csv files

-

Sampling

- A sampling script is used to to uniformly sample the thousands of files

Sampling

- A sampling script is used to to uniformly sample the thousands of files
- Because each file represents a 3 minute portion of twitter activity, there is high variance between file sizes

Sampling

- A sampling script is used to to uniformly sample the thousands of files
- Because each file represents a 3 minute portion of twitter activity, there is high variance between file sizes
- Therefore sampling must be done carefully to not distort time series data

Inspecting Data Features

- Complete json object contains ~130 attributes
- Components:
 - User data
 - Tweet data
 - Original retweet data (a copy of the outer user/tweet data)
- Reduced to 38 features in the processed data

Developing a Model Target: what to look for?

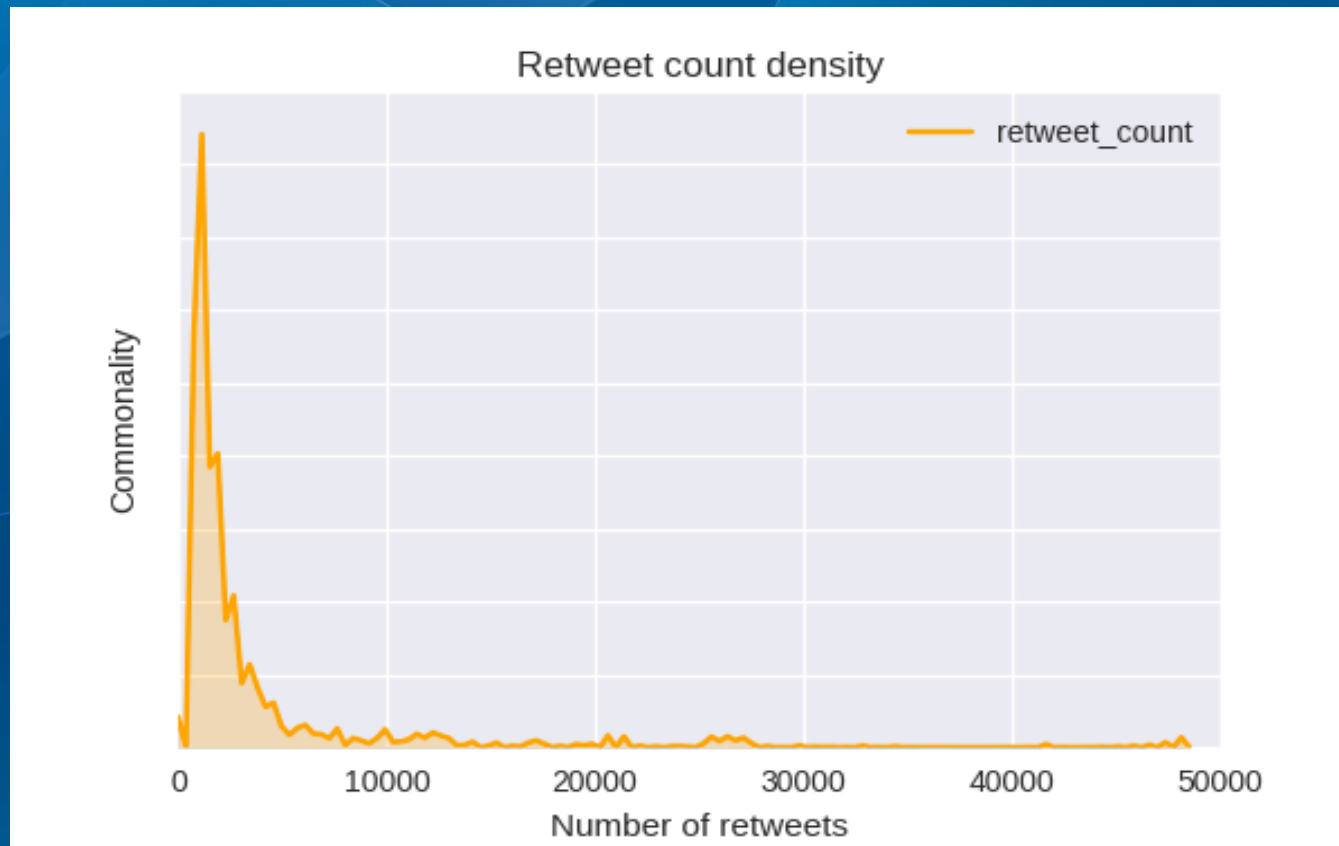


retweet_count

- The current retweet number for a given tweet

retweet_count

- The current retweet number for a given tweet
 - Not including zero retweets




Transforming retweet_count

Transforming retweet_count

- `sklearn.preprocessing.StandardScaler` applied to a log transform (offset by 1)


Transforming retweet_count

- `sklearn.preprocessing.StandardScaler` applied to a log transform (offset by 1)



Transforming retweet_count

- `sklearn.preprocessing.StandardScaler` applied to a log transform (offset by 1)



Transforming retweet_count

Doesn't work.



retweet_count is a bad target

- Very highly correlated with number of followers

retweet_count is bad

- Very highly correlated with number of followers
- Very easy to predict

retweet_count is bad

- Very highly correlated with number of followers
- Very easy to predict
- Not very interesting

Tweetability



Tweetability

- $\text{retweet_count} / \text{followers_count}$

Target: Tweetability

- $\text{retweet_count} / \text{followers_count}$
- More interesting idea for getting at the “goodness” of a tweet

Target: Tweetability

- $\text{retweet_count} / \text{followers_count}$
- More interesting idea for getting at the “goodness” of a tweet
- Foolish to presume that tweetability corresponds with “goodness”?

Target: Tweetability

- `retweet_count` / `followers_count`
- More interesting idea for getting at the “goodness” of a tweet
- Foolish to presume that tweetability corresponds with “goodness”?
 - Probably.

Target Penalty


- Add a constant penalty
 - Low retweet counts should be disproportionately rewarded for high retweet/followers ratio
 - Goal is not to maximize predictability

Tweetability

- Add a constant penalty
 - Low retweet counts should be disproportionately rewarded for high retweet/followers ratio
 - Goal is not to maximize predictability
- Add a penalty for diminishing returns


$$\frac{\text{retweetcount} + \log(\text{followerscount})}{\text{followerscount} + C / \text{retweetcount}}$$

Tranforming Tweetability



Transforming Tweetability


- After Transformations



- Minimized variance made partitioning better than k-means clustering for categorization

Transforming Tweetability

- After Transformations



Engineered features

- Hashtag count, user mention count
- sentiment polarity
- text diversity
- punctuation score
- word count
- hashtag popularity
- Etc.

Models

- RandomForest and simple decision trees quite well
- Both achieve about 70% accuracy, but high log loss
- Interestingly, the additional features past three or four core features (statuses count, favorites count, etc) barely improved accuracy (~1 point) and often harmed log loss.

Models (cont.)

- Voting classifier
 - Consists of Random forest, Naive Bayes, and Logistic regression estimators
 - Originally reduced log loss
 - Yet to run again on corrected data

Takeaways and thoughts

- Large datasets are brutal.
- Carefully examine foreign data, particularly variable json, for structure and message, not just content
 - The invested time is worth it.
- Feature engineering can be very unintuitive